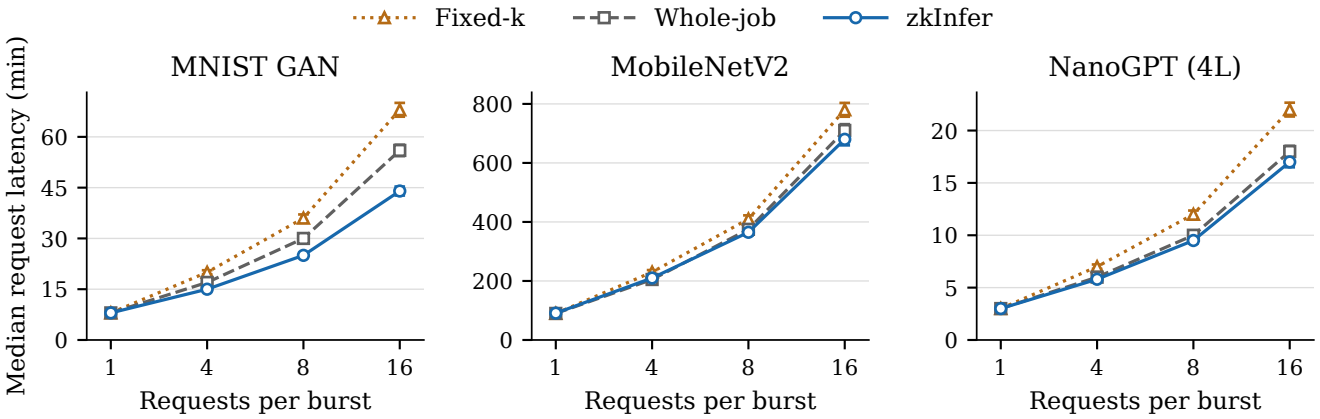
Fixed-k
Whole-job
zkInfer
MNIST GAN
MobileNetV2
NanoGPT (4L)
Median request latency (min)
Requests per burst
Requests per burst
Requests per burst
0
15
30
45
60
0
200
400
600
800
0
5
10
15
20
1
4
8
16