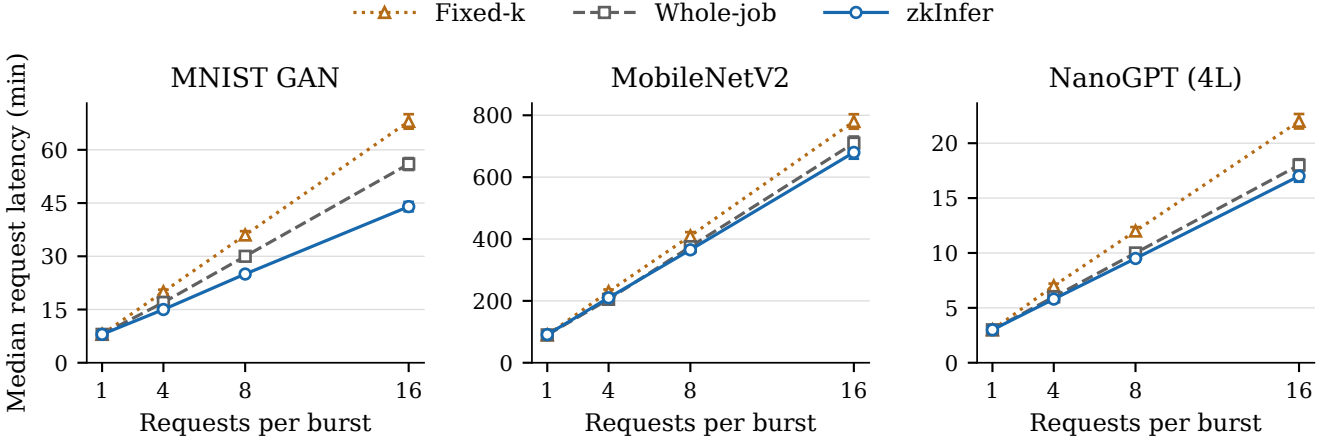
Fixed-k
Whole-job
zkInfer
MNIST GAN
MobileNetV2
NanoGPT (4L)
Median request latency (min)
0
15
30
45
60
200
400
600
800
5
10
15
20
1
4
8
16
Requests per burst
ILLUSTRATIVE DATA ONLY — NOT EXPERIMENTAL RESULTS